



Chapter 4: Intermediate SQL



Chapter 4: Intermediate SQL

- Join Expressions
- Views
- Transactions
- Integrity Constraints
- SQL Data Types and Schemas
- Authorization



Joined Relations

- A join operation
 - is a Cartesian product
 - requires that tuples in the two relations match (under some condition)
 - specifies the attributes that are present in the result of the join



Joined Relations

- **Join operations** take two relations and return as a result another relation.
- Typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

<i>Join types</i>
inner join
left outer join
right outer join
full outer join

<i>Join Conditions</i>
natural
on <predicate>
using (A_1, A_1, \dots, A_n)



Join operations – Example

□ Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

□ Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

□ Observe that

prereq information is missing for CS-315 and
course information is missing for CS-347



Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.



Left Outer Join

- course **natural left outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>



Right Outer Join

- course natural right outer join prereq

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101



Full Outer Join

□ course **natural full outer join** *prereq*

(Note: Check Latest MYSQL)

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101



Joined Relations – Example 2

- course **inner join** prereq **on**
 $course.course_id = prereq.course_id$

course_id	title	dept_name	credits	prereq_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- What is the difference between the above, and a natural join?
- course **left outer join** prereq **on**
 $course.course_id = prereq.course_id$

course_id	title	dept_name	credits	prereq_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	null	null



Joined Relations – Examples

- course **natural right outer join** prereq

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101

- course **full outer join** prereq **using** (course_id)

(Note: Check Latest MySQL)

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101



Difference between on and where

□ **select ***

```
from course left outer join prereq on  
course.course_id = prereq.course_id
```

course_id	title	dept_name	credits	prere_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	null	null

□ **select ***

```
from course left outer join prereq on true  
where course.course_id = prereq.course_id
```

course_id	title	dept_name	credits	prere_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190



Difference between on and where contd.

- Every tuple satisfies “**on** true”
- Therefore there are no dangling tuples
- **on** condition is a part of **join** clause, but **where** is not
- **where** clause is evaluated after **from** clause



Views



Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)
- Consider a person who needs to know instructors name and department, but not the salary. This person should see a relation described, in SQL, by

```
select ID, name, dept_name  
from instructor
```



Views contd.

- A **view** provides a mechanism to hide certain data from the view of certain users.
- Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**.
- Not pre-computed and stored



View Definition

- A view is defined using the **create view** statement which has the form

create view *v* as < query expression >

where <query expression> is any legal SQL expression. The view name is represented by *v*.

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.



View definition contd.

- View definition
 - is not the same as creating a new relation by evaluating the query expression
 - Rather, it causes the saving of an expression
 - The saved expression is substituted into queries using the view name.



Example Views

- A view of instructors without their salary

create view *faculty* **as**

```
select ID, name, dept_name  
from instructor
```

- Find all instructors in the Biology department

```
select name
```

```
from faculty
```

```
where dept_name = 'Biology'
```

- Create a view of department salary totals

create view *departments_total_salary*(*dept_name, total_salary*) **as**

```
select dept_name, sum (salary)
```

```
from instructor
```

```
group by dept_name;
```



Views Defined Using Other Views

- **create view physics_fall_2009 as**
select course.course_id, sec_id, building, room_number
from course, section
where course.course_id = section.course_id
and course.dept_name = 'Physics'
and section.semester = 'Fall'
and section.year = '2009';

- **create view physics_fall_2009_watson as**
select course_id, room_number
from physics_fall_2009
where building= 'Watson';



View Expansion

- Expand use of a view in a query/another view

```
create view physics_fall_2009_watson as
(select course_id, room_number
from (select course.course_id, building, room_number
      from course, section
     where course.course_id = section.course_id
       and course.dept_name = 'Physics'
       and section.semester = 'Fall'
       and section.year = '2009')
  where building= 'Watson';
```



Update of a View

- Add a new tuple to *faculty* view which we defined earlier

insert into faculty values ('30765', 'Green', 'Music');

This insertion must be represented by the insertion of the tuple

('30765', 'Green', 'Music', null)

into the *instructor* relation



Some Updates cannot be Translated Uniquely

- **create view *instructor_info* as**
select *ID*, *name*, *building*
from *instructor*, *department*
where *instructor.dept_name*= *department.dept_name*;
- **insert into *instructor_info* values ('69987', 'White', 'Taylor');**
 - ▶ which department, if multiple departments in Taylor?
 - ▶ what if no department is in Taylor?
 - ▶ if nulls are inserted for unknown values, the view *instructor_info* still does not include ('69987', 'White', 'Taylor')
- Therefore, modifications are generally not permitted on view relations, except in limited cases



Updates on views in SQL

- Most SQL implementations allow updates only on simple views
 - The **from** clause has only one database relation.
 - The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **distinct** specification.
 - Any attribute not listed in the **select** clause can be set to null
 - The query does not have a **group by** or **having** clause.



Tuples not satisfying where

- **create view** *history_instructors* **as**
select *
from *instructor*
where *dept_name*= 'History';
- What happens if we insert ('25566', 'Brown', 'Biology', 100000) into *history_instructors*?
 - SQL will allow this
- **with check option** clause at the end of view definition
 - ensures tuples not satisfying the where clause are not inserted



Materialized Views

- **Materializing a view:** create a physical table containing all the tuples in the result of the query defining the view
- If relations used in the query are updated, the materialized view result becomes out of date
 - Need to **Maintain** the view, by updating the view whenever the underlying relations are updated.
- SQL does not provide a standard way of specifying that a view is materialized



Transactions

- Unit of work
- A sequence of query and/or update statements
- Atomic transaction
 - either fully executed or rolled back as if it never occurred
- Isolation from concurrent transactions



Transactions contd.

- Transactions begin implicitly
 - Ended by **commit work** or **rollback work**
- But default on most databases: each SQL statement commits automatically
 - Can turn off auto commit for a session (e.g. using API)
 - In SQL:1999, can use: **begin atomic** **end**
 - ▶ Not supported on most databases



Integrity Constraints on a Single Relation

- **not null**
- **primary key**
- **unique**
- **check (P)**, where P is a predicate



Not Null and Unique Constraints

- **not null**

- Declare *name* and *budget* to be **not null**

name varchar(20) not null

budget numeric(12,2) not null

Primary keys do not need to be explicitly declared **not null**

- **unique** (A_1, A_2, \dots, A_m)

- The unique specification states that the attributes A_1, A_2, \dots, A_m form a candidate key.
 - Candidate keys are permitted to be null (in contrast to primary keys).



The check clause

- **check (P)**

where P is a predicate

Example: ensure that semester is one of fall, winter, spring or summer:

```
create table section (
    course_id varchar (8),
    sec_id varchar (8),
    semester varchar (6),
    year numeric (4,0),
    building varchar (15),
    room_number varchar (7),
    time_slot_id varchar (4),
    primary key (course_id, sec_id, semester, year),
    check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))
);
```



Referential integrity

- Let A be a set of attributes. Let R and S be two relations that contain attributes A and where A is the primary key of S.
 - if for any values of A appearing in R these values also appear in S, A is said to be a
 - **foreign key** of R
- In general, a referential integrity constraint does not require A to be a primary key of S



Referential integrity in SQL: direct support

- By default, a foreign key references a primary key attribute of another table

```
create table course (
    course_id  char(5) primary key,
    title       varchar(20),
    dept_name  varchar(20) references department /* foreign key */
)
```

- A list of attributes *A* of the referenced relation can be specified explicitly
 - *A* must be declared a candidate key using a **unique** constraint or a **primary key** constraint



Cascading Actions in Referential Integrity

- When a referential-integrity constraint is violated, the normal procedure is to reject the action that caused the violation.
- Alternatively

```
create table course (
```

```
...
```

```
    dept_name varchar(20),  
    foreign key (dept_name) references department  
        on delete cascade  
        on update cascade,
```

```
...
```

```
)
```

- on delete cascade** will work when you delete a tuple in *department*
- alternative actions to cascade: **set null, set default**



Integrity Constraint Violation

- E.g.

```
create table person (
    ID char(10),
    name char(40) unique,
    mother char(10),
    father char(10),
    primary key ID,
    foreign key father references person(name),
    foreign key mother references person(name))
```

- How to insert a tuple without causing constraint violation ?
 - insert father and mother of a person before inserting person
 - OR, set father and mother to null initially, update after inserting all persons (not possible if father and mother attributes declared to be **not null**)
 - OR defer constraint checking (not supported in many databases)



Assertions

- An **assertion** is a predicate expressing a condition that we wish the database to satisfy always.
- Domain and referential integrity constraints are special forms of assertions
- The following constraint can be expressed using assertions, but not using domain / referential integrity constraints:
 - For each tuple in the *student* relation, the value of the attribute *tot_cred* must equal the sum of credits of courses that the student has completed successfully.
- An assertion is tested for validity – often time consuming
- Therefore not supported in many RDBMs



Assertions contd.

- **create assertion** <assertion-name> **check** (<predicate>);
- **Example**

```
create assertion credits_earned_constraint check
not exists (select ID
            from student S
            where tot_cred <> (select sum(credits)
                                  from takes, course
                                  where takes.course_id = course.course_id
                                        and
                                        S.ID= takes.ID.and
                                        takes.grade <> 'F' and
                                        takes.grade is not null));
```



Index Creation

- Many queries reference only a small proportion of the records in a table.
- Inefficient for the system to read every record
- An **index** on an attribute of a relation is
 - a data structure that allows the database system to find those tuples in the relation that have a specified value for that attribute efficiently
 - without scanning through all the tuples of the relation.
- We create an index with the **create index** command
create index <name> on <relation-name> (attribute);



Index Creation contd.

- **create table student**
*(ID varchar (5),
name varchar (20) not null,
dept_name varchar (20),
tot_cred numeric (3,0) default 0,
primary key (ID))*
- **create index studentID_index on student(ID)**

e.g. **select ***
from student
where ID = '12345'

can be executed by using the index to find the required record, without looking at all records of *student*

More on indices later



Large-Object Types

- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*:
 - **blob**: binary large object -- uninterpreted binary data
 - **clob**: character large object -- a large collection of character data
- When a query returns a large object, a pointer is returned rather than the large object itself.



User-Defined Types

- **distinct types** in SQL
- **create type** construct in SQL creates user-defined type

```
create type Rupees as integer final;  
create type Dollars as integer final;
```

/ Check latest MySQL */*

- Example:

```
create table department  
(dept_name varchar (20),  
building varchar (15),  
budget Dollars);
```



Domains

- **create domain** construct in SQL-92 creates user-defined domain types /* Check latest MySQL */

```
create domain person_name char(20) not null
```

- Example for a constraint:

```
create domain degree_level varchar(10)
constraint degree_level_test
check (value in ('Bachelors', 'Masters', 'Doctorate'));
```

- Types and domains - differences
 - Domains can have constraints, such as **not null**
 - Domains can have **default** values for variables



Authorization

- We may assign a user several forms of authorizations on parts of the database (related to data).
 - **Read** - allows reading, but not modification of data.
 - **Insert** - allows insertion of new data, but not modification of existing data.
 - **Update** - allows modification, but not deletion of data.
 - **Delete** - allows deletion of data.
- Each of these types of authorizations is called a **privilege**.
- We may authorize the user
 - all, none, or a combination of these types of privileges
 - on specified parts of a database, such as a relation or a view.



Authorization Specification in SQL - data

- The **grant** statement is used to confer authorization
 - grant <privilege list> on <relation or view > to <user list>**
- <user list> is:
 - a user-id
 - **public**, which allows all valid users the privilege granted
 - A role (more on this later)
- Example:
 - **grant select on department to Amit, Satoshi**
- Granting a privilege on a view does not imply granting any privileges on the underlying relations.
- The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).

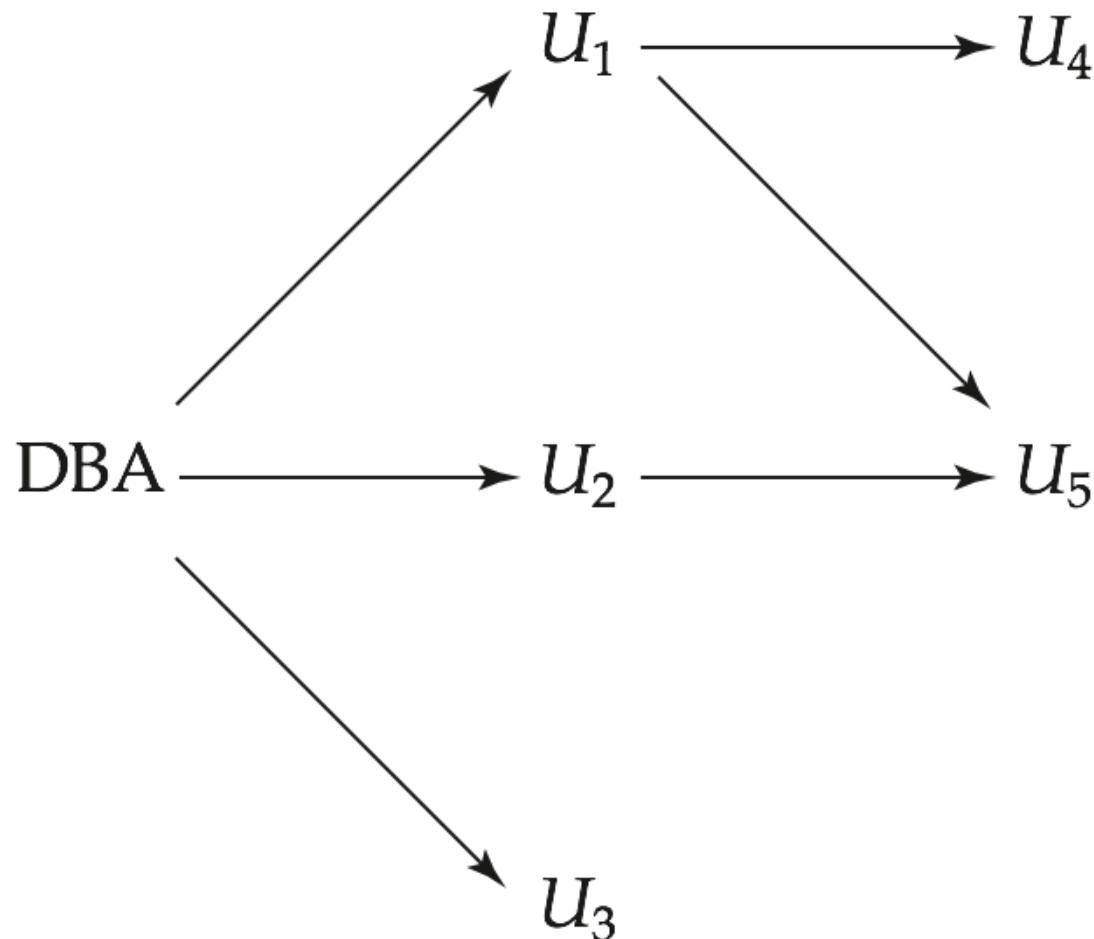


Privileges in SQL - data

- **select**: allows read access to relation, or the ability to query using the view
 - Example: grant users U_1 , U_2 , and U_3 **select** authorization on the *instructor* relation:
$$\text{grant select on instructor to } U_1, U_2, U_3$$
- **insert**: the ability to insert all or some attributes of tuples
- **update**: the ability to update all or some attributes
- **delete**: the ability to delete tuples.
- **all privileges**: used as a short form for all the allowable privileges



Authorization-Grant Graph



Transfer of privileges:
grant select on department to U_1 with grant option;



Revoking Authorization in SQL

- The **revoke** statement is used to revoke authorization.
revoke <privilege list> on <relation or view> from <user list>
- Example:
revoke select on student from U₁, U₂, U₃
- <privilege-list> may be **all** to revoke all privileges the revoker may hold.
- If <revoker-list> includes **public**, all users lose the privilege except those granted it explicitly.
- If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation.
- All privileges that depend on the privilege being revoked are also revoked – view and underlying relations



Roles

- A **role** is a way to distinguish among various users
 - denotes what these users can access/update in the database.
 - No need to give individual authorizations
- To create a role we use:
create role <name>
- Example:
 - **create role instructor**
- Once a role is created we can assign “users” to the role using:
 - **grant <role> to <users>**



Roles Example

- **create role** instructor;
- **grant** *instructor* **to** Amit;
- Privileges can be granted to roles:
 - **grant select on** *takes* **to** *instructor*;
- Roles can be granted to users, as well as to other roles
 - **create role** teaching_assistant
 - **grant** *teaching_assistant* **to** *instructor*,
 - ▶ *instructor* inherits all privileges of *teaching_assistant*
- Chain of roles
 - **create role** dean;
 - **grant** *instructor* **to** *dean*;
 - **grant** *dean* **to** Satoshi;



Authorization on Views

- **create view geo_instructor as**
(select *
from instructor
where dept_name = 'Geology');
- **grant select on geo_instructor to geo_staff**
--geo_staff is a role
- Suppose that a geo_staff member issues
 - **select * from geo_instructor;**
- What if
 - geo_staff does not have permissions on *instructor*?
 - creator of view did not have some permissions on *instructor*?



Other Authorization Features

- transfer of privileges
 - **grant select on department to Amit with grant option;**
 - **revoke select on department from Amit, Satoshi cascade;**
 - ▶ Revocation will cascade (default behavior)
 - **revoke select on department from Amit, Satoshi restrict;**
 - ▶ An error if there are cascading revocations



Authorizations on schema

- Primitive: Only the owner can carry out any modification to the schema
- **references** privilege to declare foreign keys while creating a relation
 - **grant reference** (*dept_name*) **on** *department* **to** Mariano;
 - Now Mariano can create a foreign key in *r* referencing *department*
 - why is this a privileged operation ?



End of Chapter 4